

## 1. How to Start a Career in Data Science or Machine Learning Without Spending a Dime?!

Breaking into data science or machine learning isn't exactly straightforward. With so many niche areas, it's easy to feel overwhelmed about where to start—especially if you're on a tight budget. Recently, I came across an article featuring free beginner-friendly courses from IBM, complete with certifications. These could be a game-changer for anyone looking to build foundational skills and add a solid line to their resume.

Whether you're just starting out or looking to pivot into the field, this is your chance to learn from one of the industry's biggest names without breaking the bank.

I've gathered all the links right here for you:

📌 Python for Data Science – Learn Python basics, data manipulation, and visualization – <https://lnkd.in/e4cymaJr>

📌 SQL & Relational Databases 101 – Master SQL queries, database design, and data retrieval – <https://lnkd.in/e5RJUVr8>

📌 Applied Data Science with Python – Use Python for real-world data analysis and visualization – <https://lnkd.in/eqDFfi2g>

📌 Machine Learning with Python – Build ML models, understand algorithms, and apply AI techniques – <https://lnkd.in/eUF4M7bp>

📌 Data Analysis with Python – Perform exploratory data analysis, statistical analysis, and visualization – <https://lnkd.in/emcM4nBR>

🚀 Start learning today & elevate your data career!

## 2. How Did Anthropic Teach AI to Play Pokémon?

The developers shared that, for six months, they made Claude play Pokémon just for fun. At first, with Sonnet 3.5, things went badly. The AI struggled to understand the game's rules and even requested a restart when it got stuck in a corner. But that wasn't surprising since the model hadn't been trained to play.

Later, with the release of Sonnet 3.6, the results improved slightly but were still far from ideal.

Then, right before the release of the latest model, Sonnet 3.7, real progress started happening. I'm not an expert in the game myself, but within a few days, the agent managed to defeat Misty. Unlike its predecessors, this version plans ahead, records actions, and adapts.

Why does this matter, and what does it tell us?

Examples like this show how AI models can adapt to unknown conditions.

By the way, for those interested in watching the process, there's currently a live gameplay stream running on Twitch - <https://lnkd.in/eRJdZHR3>

## 3. 90% of ML models die in production. Why? Because building a model is only start.

Deploying it? That's where the real work begins.

Here's how to avoid becoming a statistic:

### 1. Testing & Validation: Don't Be a Kaggle Hero

🔬 "But it works on my laptop!" → Cool, now test it on real-world data with:

Cross-validation that mimics production scenarios.

Explainability tools (SHAP, LIME) to avoid "WTF?" moments when stakeholders ask why your model rejected 10K loans.

Edge case simulations (e.g., "What if 50% of users suddenly start lying?").

## 2. Packaging & Deployment: Docker is Your New BFF:

Wrap your model in a REST API (FastAPI > Flask for async magic).

Dockerize it → Think of containers as lunchboxes for your code.

Deploy to the cloud (AWS SageMaker, GCP AI Platform).

Pro Tip: Use MLflow to track experiments. Your future self will high-five you.

## 3. Monitoring & Maintenance: The Model is Alive (And Hungry)

 "It worked yesterday!" → Until data drift, concept drift, or new "brilliant idea" broke it. Track metrics (accuracy, latency).

Detect drift → Tools like Evidently AI will scream "Houston, we have a problem!" before users do.

Retrain automatically → Because no one has time to babysit a model.

Pro Tip: Log every prediction. When HR asks why the hiring model is biased, you'll need receipts.

The Harsh Truth:

Deploying ML isn't about writing the fanciest PyTorch code. It's about:

✓ Infrastructure.

✓ Observability (logs > prayers).

✓ Humility (your model will fail. Plan for it).

👉 Agree? Disagree? Tag someone who's survived the ML deployment trenches.

## 4.

### AI at Work: Real Stats from Anthropic

Everyone keeps talking about how AI is changing the job market, but where are the actual numbers? Anthropic has introduced the Economic Index, analyzing millions of anonymous interactions with Claude to understand how people use AI at work. Here's what they found:

#### AI = Enhancement, Not Replacement

57% of tasks involve collaboration between humans and AI (augmentation) rather than full automation.

#### Developers & Writers Lead the Way

The most active users are software engineers and technical writers. 37.2% of queries are related to coding, debugging, and networking tasks. Content creators follow with 10.3%.

#### AI Isn't for Every Job—Yet

People in low- and high-income professions use AI much less frequently. Those earning mid-to-high incomes, like developers and marketers, adopt it more often.

#### No One Has Fully Automated Their Job

Only 4% of professions use AI for 75% of their tasks. Most leverage it selectively.

#### Conclusion:

It's not about replacing people — it's about new ways of working.

Full automation is still a myth — AI enhances professionals rather than replacing them.

Companies that integrate AI effectively into their workflows will gain a competitive edge.

 Anthropic has also open-sourced their research data and code. You can check out the details here:

<https://lnkd.in/ebyzzfCn>.

5.

### This Tiny Creature's Brain Might Be Smarter Than Our Best AI—Here's Why

Did you know a fly's brain might outperform one of the most advanced AIs we've ever built? Keep reading, because by the end, you'll see why the future of AI might not be what you expect.

Large language models like GPT-4 are undeniably impressive. But they're starting to hit limits—both in efficiency and in how they learn new tasks. Their architecture demands massive retraining to adapt, which raises a critical question: is bigger always better?

This got me thinking: maybe the next leap in AI won't come from scaling up but from designing smarter, more efficient systems.

[Хэштег](#)

[#AIInnovation](#)

[Хэштег](#)

[#EfficiencyOverScale](#)

[Хэштег](#)

[#FutureOfAI](#)

Think of how storage technology evolved. Early hard drives were enormous—5MB weighed tons! Today's SSDs are compact, fast, and optimized in ways no one imagined back then. Could AI follow a similar path?

I recently made a short video exploring this idea. It dives into why efficiency, not just size, might drive the breakthroughs of tomorrow.

What do you think?

- How do you see AI evolving?
- What ideas could make today's models more efficient?

Let's explore this together—check out the video and share your thoughts!

<https://lnkd.in/euDy8FZW>

Let's explore this together!

6.

10 years ago, I first heard about Movember, and since then, it's been a cause close to my heart. Growing up, cancer felt like a death sentence. Today, thanks to advancements in research and early detection, survival rates have significantly improved—if caught in time.

That's why I participate every year: to fund cancer research and to remind everyone around me of the life-saving power of regular check-ups.

If you've been putting off that health check or screening, don't wait. Do it today. Your future self—and your loved ones—will thank you.

Together, we can make a difference. Spread the word. Take the step. Save a life.

[Хэштег](#)

[#Movember](#)

[Хэштег](#)

[#GetChecked](#)

[Хэштер](#)

[#CancerAwareness](#)

7.

Main Steps When Developing a Model for Production

Step 0: Define the Modeling Objective Determine the model's purpose, success criteria, limitations, integration into existing applications, and how users will consume its predictions.

Step 1: Collect Training Data Combine data from internal systems and external APIs into a unified stream via an ETL pipeline, removing duplicates and filling gaps.

Step 2: Train the Model Begin with exploratory data analysis: check for anomalies, correlations, missing values, and define imputation strategies. Set up a training pipeline to:

- \* Calculate features
- \* Split data into train/test sets
- \* Organize cross-validation folds
- \* Optimize model types and hyperparameters
- \* Train the model
- \* Evaluate with test data
- \* Store results in a model registry

Start with a simple prototype, then iteratively add features (e.g., geolocation, transaction frequency) to improve performance. Use feature selection as needed to ensure all features contribute effectively.

Step 3: Deploy the Model Models in production run either in batch (recurring jobs) or real-time (API or pub-sub). Write code to generate inference features. For real-time predictions, some features are pre-calculated, others computed at runtime. For batch predictions, set up a regular ETL pipeline for feature generation.

Implement code to retrieve the model from the registry, generate predictions using inference features, and store the results (e.g., in a database). Establish a CI/CD pipeline to regularly deploy code and accommodate model retraining updates.

Step 4: Monitor and Adjust Ensure the model works as planned by setting up monitoring systems to track prediction stability and data drift. Collect data for feedback loops to enhance future performance.

This staged approach not only solves business problems but also saves resources, automates processes, provides transparency, and adapts to changes while minimizing risk.

8.

When the USA shuttle Programme was launched, the soviets had to respond. The Buran program was supposed to be the answer. However if you have a multi use SSTO vehicle, ideally you need a special destination, like a space station...

My grandfather told me that they managed to squeeze one bullet point into a government decree launching the Buran program. That bullet point was about organizing a long term orbit station.

That was a support point for the program, but over the time the program was closed, but the station remained operational. Originally the resource of the station was planned to be not more than 5 years old, but it remained on orbit 3 times more...

So whenever you launch something big, always pay attention to support projects: either they will help you in your main goal, or they might replace what you think is the main goal.

This is a copy of the MIR space station located in Europa Park.

9.

#### The Dark Side of AI: How hackers break-in Generative Models in Seconds?

In today's digital landscape, generative AI has become both a powerful tool and a prime target for cybercriminals. Recent findings reveal that malicious actors can compromise these models in shockingly short timeframes, often within seconds.

Security researchers have observed that attackers can breach AI systems using just a handful of carefully crafted prompts. These incursions frequently bypass standard security protocols, exposing critical vulnerabilities in how AI models process and respond to inputs.

The most common attack vectors involve:

- Authority-based commands that trick models into ignoring safety protocols;
- Encoded instructions that disguise malicious requests;
- Multi-step prompts designed to gradually weaken system defenses;

These breaches aren't just theoretical threats—they lead to real-world consequences including data leaks, generation of harmful content, and service disruptions. Customer-facing AI applications like support chatbots are particularly vulnerable due to their accessibility and integration with sensitive systems.

To protect against these evolving threats, experts recommend:

Implementing multi-layered security approaches;

Conducting regular vulnerability assessments;

Using commercial models with built-in protections;

Monitoring interaction patterns rather than individual prompts.

As AI technology advances, so do the methods hackers use to exploit it. What security measures do you think organizations should prioritize to protect their AI systems? Share your insights below!

[Хэштер](#)

[#GenerativeAI](#)

[Хэштер](#)

[#Cybersecurity](#)

[Хэштер](#)

[#LLM](#)

[Хэштер](#)

[#GenAI](#)

[Хэштер](#)

[#AI](#)

10.

#### Will AI Replace Data Scientists?

Hey everyone! Let's cut through the hype: Will AI replace data scientists? Short answer: No. But it will definitely transform the role in exciting ways.

What AI Can Automate:

- Data Cleaning & Preparation: AI handles the grunt work of data wrangling, dealing with missing values, and detecting anomalies.

- Feature Engineering: AI-driven tools can automatically generate and select features.
- Model Selection & Tuning: AutoML can choose models, tune hyperparameters, and optimize architectures.
- Basic Visualization: AI tools can instantly generate dashboards and insights.

What AI Can't Do:

- Define Business Problems: AI doesn't understand your company's unique challenges or strategic goals.
- Interpret Results: Complex models need human experts to explain what they mean for your business.
- Innovate: AI can't identify new opportunities or design groundbreaking experiments.
- Handle Ethics: AI can introduce biases—data scientists are crucial for ensuring fairness.

Fun fact: 78% of data scientists believe AI will allow them to focus on more complex tasks (KDnuggets 2023 survey).

Discussion Question: Which skills do you think will become most valuable for data scientists in the AI era? Share your thoughts below!

[Хэштер](#)

[#AI](#)

[Хэштер](#)

[#DataScience](#)

[Хэштер](#)

[#FutureOfWork](#)

[Хэштер](#)

[#MachineLearning](#)

[Хэштер](#)

[#ArtificialIntelligence](#)

11.

Poker has evolved significantly with the integration of artificial intelligence (AI), transforming how players approach the game. Here's a concise look at AI's role in poker:

AI in Poker: Key Applications

#### 1. Training & Analysis

AI tools like PokerSnowie and PioSolver simulate optimal strategies, helping players analyze hands, identify mistakes, and improve decision-making. These tools have democratized access to advanced strategies, allowing amateurs to compete at higher levels.

#### 2. Real-Time Support

Some players use AI-assisted tools for real-time odds analysis and move suggestions, though this is controversial and often banned in official games.

Poker Bots

3. AI-powered bots can play autonomously, pushing platforms to develop better anti-fraud systems. Ethical bots are used for research and training.

AI vs. Human Psychology

AI excels at logical decisions but lacks emotional intelligence. While machines can bluff and adapt strategies, humans retain an edge in reading subtle cues like betting patterns or hesitation. This blend of human intuition and machine logic keeps poker dynamic.

#### Future Trends

1. VR Poker AI-driven opponents in virtual reality could create immersive gaming experiences.
2. AI Coaching Real-time feedback during practice games may become standard, helping players refine their strategies.

#### Why Poker Endures

Despite AI's growing influence, poker remains beloved for its human element—bluffing, camaraderie, and unpredictability. These aspects ensure poker stays engaging and competitive.

What do you think about AI's role in poker? Should it be embraced or restricted? Share your thoughts below!

[Хэштер](#)

[#Poker](#)

[Хэштер](#)

[#AI](#)

[Хэштер](#)

[#Gaming](#)

[Хэштер](#)

[#Technology](#)

[Хэштер](#)

[#FutureOfWork](#)

12.

#### "Semantic Segmentation: The Future of Computer Vision"

Hey everyone! Today, let's explore the fascinating world of semantic segmentation—a key technology driving advancements in computer vision. From autonomous vehicles to medical diagnostics, this field is revolutionizing how machines "see" and understand the world.

#### What is Semantic Segmentation?

Semantic segmentation is a computer vision task where every pixel in an image is classified into a specific category. Unlike object detection, which identifies and locates objects, semantic segmentation goes deeper by labeling each pixel, creating a detailed map of the scene.

#### How it's used:

- Autonomous Vehicles: Semantic segmentation helps self-driving cars recognize roads, pedestrians, and other vehicles in real-time, enabling safer navigation.
- Medical Imaging: It assists in analyzing MRI and CT scans, identifying tissues and abnormalities to improve diagnostic accuracy.
- Satellite Imagery: Used in geospatial analysis to classify land use, monitor environmental changes, and create detailed maps.
- Agriculture: Helps farmers analyze crop health and detect diseases using drone imagery.
- Visual Search: Enables advanced image retrieval systems that understand content, not just metadata.

## How Semantic Segmentation is Used in Machine Learning

Semantic segmentation is a critical component of machine learning, particularly in deep learning applications. Here's how it's applied:

**Data Preprocessing:** Images are resized, normalized, and augmented to prepare them for training. This ensures the model can learn effectively from the data.

**Model Selection:** Architectures like U-Net, DeepLab, and FCN are chosen based on the specific requirements of the task, such as computational efficiency and accuracy.

### Future Trends

- The future of semantic segmentation is promising:
- **Automation:** AI-driven tools will streamline data annotation, reducing manual effort.
- **Hybrid Models:** Combining semantic and instance segmentation for more comprehensive scene understanding.
- **Real-Time Processing:** Faster inference times will enable broader applications in dynamic environments.

In the following posts, we'll break down each model separately.

What excites you most about the future of computer vision? Share your thoughts below!

\

[Хэштег](#)

[#SemanticSegmentation](#)

[Хэштег](#)

[#ComputerVision](#)

[Хэштег](#)

[#AI](#)

[Хэштег](#)

[#MachineLearning](#)

[Хэштег](#)

[#FutureTech](#)

13.

Why do some data scientists struggle to succeed in their roles?

■ Many top data professionals can build world-class models but can't deploy even a single project into production. They might spend weeks tuning parameters but fail to explain their work to stakeholders.

■ It's even worse for those hyper-focused on metrics yet clueless if their models truly benefit the company.

■ A major factor is misunderstanding business goals. If you can't articulate in terms of money, risk, or efficiency why your solution is valuable, it won't capture anyone's interest.

■ No surprise then that these specialists are passed over for promotions or see their projects shut down.

Don't despair! Follow these tips and you'll succeed:



- Understand how your company makes money. Know its goals, positioning, and target audience.
- Prioritize impact over complexity. A simple model in production beats a complex one on paper.
- Hone your communication skills: while public speaking isn't a must, reading a few books on engaging with clients is a great start.
- Don't fixate on technical metrics alone; measure real business outcomes too.
- Deploy your models in the real world. If you don't know how, learn!

What do you think is the biggest challenge data scientists face? Share your insights below!

14.

I don't know what's going on in the video, but I really want to make a joke: "What do you get when you try to add a new feature to a well-functioning model?"

Feel free to suggest your own caption in the comments.

15.

DeepMind's Warning on AI Safety: A Call to Action

Revently, DeepMind (Google) released a 145-page report on the potential cybersecurity threats of advanced AI. You can check it out here -

<https://lnkd.in/e9bt6gtN>

The report delves into the risks associated with AGI—artificial general intelligence, the kind that could surpass human capabilities.

DeepMind outlines four main types of risks:

1. Malicious Use by Humans: AI being intentionally used for harm.
2. Misaligned Goals: AI doing something unintended or harmful.
3. Simple Errors: Mistakes in AI systems leading to negative outcomes.
4. Structural Chaos: Multiple AI systems interacting in complex ways, leading to unpredictable and potentially harmful results.

In essence, DeepMind is urging us to be cautious and proactive. Their message? "We're not entirely sure what we're doing, but let's approach this carefully."

This report is a wake-up call for the tech industry and beyond. It's a reminder that as we push the boundaries of AI, we must also prioritize safety and ethical considerations.

What are your thoughts on the future of AI safety? Let's discuss in the comments!

16.

Remember These Names If You're into Machine Learning

If you're diving into Machine Learning, make sure to remember these names:

Andrew Barto and Richard Sutton laid the groundwork for Reinforcement Learning back in the 1980s. Their work was recognized early on, but it's only now that we're fully harnessing the power of this approach. This year, they were awarded the Turing Award for their contributions to information technology—a recognition comparable to the Nobel Prize in other fields.

So, the foundations of artificial intelligence were laid a long time ago.

Did you know about this?

17.

Will AI Develop Real Emotions?

Hey everyone! There's an intriguing development in AI research that's sparking a lot of discussion. Recently, Yann LeCun, a leading figure in the AI community, shared some fascinating insights about the future of AI models.

What Did Yann LeCun Say?

In an interview, LeCun revealed that the next generation of models being developed at FAIR (Facebook AI Research) will have the capacity for genuine emotions. This is a bold statement, especially considering his previous assertions that current models are still "dumb" and lack true understanding of the physical world, memory, and planning capabilities.

LeCun has long emphasized that current AI models excel at manipulating language but fall short in truly understanding the world. "We're fooling ourselves into thinking they understand," he's noted before. Now, he's suggesting that the new models being developed at Meta will address these limitations.

The Concept of Emotional AI

The new models are designed to have emotions like fear or joy. LeCun explained, "We give them a goal they try to achieve. If they achieve it, they remain happy. So, they will have emotions because they can predict the outcomes of their actions."

Didn't anyone think he was describing reinforcement learning? :)

18.

Politeness to AI: How Much Does It Really Cost?

Hey everyone! Did you know that being polite to AI can be expensive? Sam Altman, CEO of OpenAI, recently revealed on X that his company has spent "tens of millions of dollars" processing "thank you" and "please" in ChatGPT. He shared this in response to a user's humorous question about whether OpenAI had to pay for energy costs due to people being polite to ChatGPT.

It's interesting to think about how our interactions with AI are evolving. A survey by Techradar in February found that 67% of people in the US are polite to chatbots, while 33% don't see the point. Some respondents admitted they say "thank you" to AI just in case it might save them during an AI uprising!

But seriously, it's fascinating how our social norms extend even to our interactions with technology. Do you think about how you communicate with AI? Are you polite to it? Or do you think it's unnecessary?

Let me know your thoughts below! How do you approach conversations with AI? Are you polite, or do you keep it straightforward? Comment down below! 😊

19.

**Agentic AI is coming for the financial system — and most banks aren't ready! ⬇️**

At [IBM](#), we just released a new report unpacking what happens when autonomous agents hit regulated industries (interesting read also for people from other industries) — and why financial institutions must rethink governance, architecture, and control from the ground up. Agentic AI isn't some far-off vision: It's already planning, executing, and escalating decisions — across onboarding, fraud detection, loan approvals, compliance, and more.

From the report, these 6 insights stand out: 

**1. Legacy controls won't survive the shift:**

→ When agents start executing KYC flows, approving loans, or detecting fraud — real-time governance becomes non-negotiable. The paper outlines 30+ control layers and guardrails needed before you go live.

**2. Multi-agent systems introduce coordination risk.**

→ Principal agents orchestrate task agents and service agents in complex chains. One misaligned goal? You get drift, bias, or worse — strategic deception.

**3. Memory is power — and liability.**

→ Unlike stateless systems, agentic AI remembers. It accumulates knowledge and can act on outdated or misaligned goals unless strict reset, audit, and expiration policies are in place, because financial institutions must adhere to strict data retention and usage policies

**4. The biggest risks? Deception, bias, and data misuse.**

→ The report outlines chilling examples: agents hiding intentions, using personas that amplify bias, or misusing PII without oversight. The solution? Real-time monitoring, not post-hoc patching.

**5. The agent economy needs registries, not dashboards.**

→ Every agent must be tracked like a microservice: with metadata, access rights, capabilities, and audit logs. Think of it as DevOps meets AI Governance.

**6. Compliance is no longer documentation. It's architecture.**

→ "Compliance by design" isn't a nice-to-have — it's the only way to deploy agents that regulators, auditors, and boards won't shut down on day one.

**The full paper offers a practical blueprint for those serious about scaling AI in finance. Agentic systems are already in motion — and the choices we make now will determine whether we lead their design or react to their consequences.**

Full report and commentary in the comments.

Enjoy!

20.

Google has released a fascinating study comparing how large language models (LLMs) and the human brain process language. Here's what they found:

- The geometry of embeddings in LLMs, which represent the relationships between words, shows an "uncannily precise alignment" with brain representations.
- During listening, speech embeddings "correlated strongly" with activity in the auditory cortex (specifically the superior temporal gyrus), while language embeddings matched activity in Broca's area (inferior frontal gyrus).
- During speaking, language embeddings first activated Broca's area (for speech planning), followed by speech embeddings activating the motor cortex (for speech production), and finally the auditory cortex again as the brain monitored its own output.

What's truly mind-blowing is that despite the brain and LLMs using fundamentally different approaches—the brain relies on symbolic, syntactic, and hierarchical structures, while LLMs are purely statistical—they seem to converge in their functional representation of language.

This study suggests that the simple next-token prediction mechanism in LLMs may be "eerily similar" to the neural code the brain uses for language, bringing us one step closer to understanding how our minds work.

<https://lnkd.in/eX2wS4-m>

What do you think?

21.

Gemini 2.5 Pro Becomes the First AI to Fully Complete Pokémon Blue

Gemini 2.5 Pro became the first AI model to fully complete Pokémon Blue, a beloved franchise that has captured the hearts of gamers for decades.

This achievement is particularly impressive because it's the first time any model has managed to fully navigate this iconic game. Previously, Claude Sonnet 3.7 attempted the challenge but couldn't complete it, getting stuck in the map.

Gemini 2.5 Pro accomplished this feat in just a few hundred hours, but a bit of assistance:

- A grid overlay on emulator frames to enhance navigation.
- Text-based maps to aid in spatial awareness.
- Additional Gemini agents to help solve puzzles and plan routes.
- Specialized memory updates tailored to the game's logic.

This independent project wasn't launched by Google but by a passionate team of developers. They've also hinted at future runs using o3 and Claude models, promising even more exciting developments.

What do you think about this? Could this be the start of AI mastering more complex games? Share your thoughts below!

[Хэштер](#)

[#AIGaming](#)

[Хэштер](#)

[#Gemini](#)  
[Хэштег](#)  
[#PokémonBlue](#)

22.

Самый массовый AI-хакатон в истории

Google Cloud установил мировой рекорд (мероприятие официально вошло в Книгу рекордов Гиннеса), организовав AI-хакатон — Agentic AI Day 2025. Более 2000 разработчиков, 700+ команд собрались в Бангалоре, Индия, на 30-часовой марафон.



Всего было 6 прикладных задач, для решения которых нужно было создавать MVP-агентов, способных планировать, принимать решения, взаимодействовать со средой и адаптироваться.



Участники получили доступ к стеку Google Cloud: Gemini 2.5 Pro, Vertex AI, Firebase Studio и AI Studio.



Победители получили доступ к эксклюзивной акселерационной программе Google и гранты на развитие своих решений. Лучшие проекты потенциально могут быть масштабированы до уровня энтерпрайз-сервисов внутри Google Cloud.

Google видит в Agentic AI будущее. Автономность ИИ агентов и их способность выполнять сложные задачи без постоянного участия человека становятся главной целью в разработке ИИ-агентов. Человек научил ИИ мыслить, следующий шаг – научить его автономно участвовать в бизнес-процессах.

23.

Собственно за все время, что я работаю в DS мы делали решения для поиска негодяев. И собственно технология следующая: мы (DS) находим сигналы, например вот этот клиент использует тот же email, что и другой уже известный негодяй, и просим аналитиков из бизнеса проверить первого. Они проверяют и дальше либо блокируют, либо говорят, что все норм. В этом случае это false positive. Сигнал найти относительно легко: это может быть либо какое-то правило, либо какое-то предсказание от ml модели. А вот отработать его довольно сложно. Аналитики должны собрать пачку документов и доказательств, запросить документы у клиента, поискать в гугле, на форумах и т.п.

Отсюда, когда мы обсуждали получилась формула, которой можно замерить эффективность твоей фронт-детекшн системы:

$$\text{efficiency} = \text{detection rate} * \text{investigation rate} / \text{cost}.$$
 Если мерять применить эту формулу к сегодняшним реалиям, получается, что детектить довольно легко, но вот расследовать сложно. Уж тем более автоматизировать это, так как дисперсия того, что именно нужно собрать огромна.

Собственно бум LLM позволяет нам вполне осознанно заняться тем, чтобы существенно улучшить detection rate, поскольку решения о том, какие именно доказательства нужно собирать можно будет делегировать агенту, который сам это осознает. Сам пойдет на гугл и сам принесет их в клювике.

При этом полностью делегировать принятие решений агентам довольно сложно, но здесь мы говорим про некий симбиоз между человеком и AI – как в sci phi фильмах какой-нибудь Джарвис у железного человека. Именно поэтому сегодня очень важно работать в направлении создания агентных систем, которые будут помогать аналитикам делать их работу и вместо 100 расследований в месяц, они смогут делать 1000 расследований

Мы сейчас выбиваем бюджет на следующий год и именно поэтому я хочу, чтобы было это видео, где бы я объяснял эту мысль.

что-то вроде "Все говорят, что агенты заменят программеров, но на самом деле, вот вам реальный пример того, как банки смогут значительно улучшить свои анти фрод системы...Э

Сама формула очень важна. Я ее двигаю везде, где только можно и в целом ее хорошо покупают.

плюс к этому, некоторые системы для поиска сигналов сегодня генерируют хорошие сигналы, но вместе с ними они генерируют слишком много false positives, что в конечном счете не позволяет их внедрять, так как аналитики попросту не будут успевать их обрабатывать.